

```
In [ ]: ## Load libraries
import pandas as pd
import numpy as np
import sys
import matplotlib.pyplot as plt
import matplotlib.cm as cm
from keras.datasets import mnist
plt.style.use('dark_background')
%matplotlib inline
```

WARNING:tensorflow:From c:\Users\vp140\.conda\envs\pycaretenv\lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse\_softmax\_cross\_entropy is deprecated. Please use tf.compat.v1.losses.sparse\_softmax\_cross\_entropy instead.

```
In [ ]: np.set_printoptions(precision=2)
```

```
In [ ]: import tensorflow as tf
```

```
In [ ]: tf.__version__
```

```
Out[ ]: '2.15.0'
```

---

Load MNIST Data

---

```
In [ ]: ## Load MNIST data
(X_train, y_train), (X_test, y_test) = mnist.load_data()
X_train = X_train.transpose(1, 2, 0)
X_test = X_test.transpose(1, 2, 0)
X_train = X_train.reshape(X_train.shape[0]*X_train.shape[1], X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0]*X_test.shape[1], X_test.shape[2])

num_labels = len(np.unique(y_train))
num_features = X_train.shape[0]
num_samples = X_train.shape[1]

# One-hot encode class labels
Y_train = tf.keras.utils.to_categorical(y_train).T
Y_test = tf.keras.utils.to_categorical(y_test).T

# Normalize the samples (images)
xmax = np.amax(X_train)
xmin = np.amin(X_train)
X_train = (X_train - xmin) / (xmax - xmin) # all train features turn into a numb
X_test = (X_test - xmin)/(xmax - xmin)

print('MNIST set')
print('-----')
print('Number of training samples = %d'%(num_samples))
print('Number of features = %d'%(num_features))
print('Number of output labels = %d'%(num_labels))
```

MNIST set

-----

Number of training samples = 60000

Number of features = 784

Number of output labels = 10

---

A generic layer class with forward and backward methods

---

```
In [ ]: class Layer:
        def __init__(self):
            self.input = None
            self.output = None

        def forward(self, input):
            pass

        def backward(self, output_gradient, learning_rate):
            pass
```

---

CCE loss and its gradient

---

```
In [ ]: ## Define the Loss function and its gradient
        def cce(Y, Yhat):
            return(np.mean(np.sum(-Y*np.log(Yhat), axis = 0)))
            #TensorFlow in-built function for categorical crossentropy loss
            #cce = tf.keras.losses.CategoricalCrossentropy()
            #return(cce(Y, Yhat).numpy())

        def cce_gradient(Y, Yhat):
            return(-Y/Yhat)
```

---

Generic activation layer class

---

```
In [ ]: class Activation(Layer):
        def __init__(self, activation, activation_gradient):
            self.activation = activation
            self.activation_gradient = activation_gradient

        def forward(self, input):
            self.input = input
            self.output = self.activation(self.input)
            return(self.output)

        def backward(self, output_gradient, learning_rate = None):
            return(output_gradient[:-1, :] * self.activation_gradient(self.input))
```

---

Specific activation layer classes

---

```
In [ ]: class Sigmoid(Activation):
    def __init__(self):
        def sigmoid(z):
            return 1 / (1 + np.exp(-z))

        def sigmoid_gradient(z):
            a = sigmoid(z)
            return a * (1 - a)

        super().__init__(sigmoid, sigmoid_gradient)

class Tanh(Activation):
    def __init__(self):
        def tanh(z):
            return np.tanh(z)

        def tanh_gradient(z):
            return 1 - np.tanh(z) ** 2

        super().__init__(tanh, tanh_gradient)

class ReLU(Activation):
    def __init__(self):
        def relu(z):
            return z * (z > 0)

        def relu_gradient(z):
            return 1. * (z > 0)

        super().__init__(relu, relu_gradient)
```

---

Softmax activation layer class

---

```
In [ ]: ## Softmax activation layer class
class Softmax(Layer):
    def forward(self, input):
        self.output = tf.nn.softmax(input, axis = 0).numpy()

    def backward(self, output_gradient, learning_rate = None):
        ## Following is the inefficient way of calculating the backward gradient
        softmax_gradient = np.empty((self.output.shape[0], output_gradient.shape[1]))
        for b in range(softmax_gradient.shape[1]):
            softmax_gradient[:, b] = np.dot((np.identity(self.output.shape[0]) - np.atleast_2d(self.output) * self.output.T), output_gradient[:, b])

        # Return gradient w.r.t. input for backward propagation
        return(softmax_gradient)

        ## Following is the efficient of calculating the backward gradient
        #T = np.transpose(np.identity(self.output.shape[0]) - np.atleast_2d(self.output) * self.output.T)
        #return(np.einsum('ijk, ik -> jk', T, output_gradient))
```

---

Dense layer class

---

```
In [ ]: ## Dense layer class
class Dense(Layer):
    def __init__(self, input_size, output_size):
        self.weights = 0.01*np.random.randn(output_size, input_size+1) # bias tr
        self.weights[:, -1] = 0.01 # set all bias values to the same nonzero con

    def forward(self, input):
        self.input = np.vstack([input, np.ones((1, input.shape[1]))]) # bias tri
        self.output = np.dot(self.weights, self.input)

    def backward(self, output_gradient, learning_rate):
        ## Following is the inefficient way of calculating the gradient w.r.t. w
        weights_gradient = np.zeros((self.output.shape[0], self.input.shape[0]),
        for b in range(output_gradient.shape[1]):
            weights_gradient += np.dot(output_gradient[:, b].reshape(-1, 1), self.
            weights_gradient = (1/output_gradient.shape[1])*weights_gradient

        ## Following is the efficient way of calculating the weightsgradient
        #weights_gradient = (1/output_gradient.shape[1])*np.dot(np.atleast_2d(ou

        # Gradient w.r.t. the input
        input_gradient = np.dot(self.weights.T, output_gradient)

        # Update weights using gradient descent step
        self.weights = self.weights + learning_rate * (-weights_gradient)

        # Return gradient w.r.t. input for backward propagation
        return(input_gradient)
```

---

Function to generate sample indices for batch processing according to batch size

---

```
In [ ]: ## Function to generate sample indices for batch processing according to batch s
def generate_batch_indices(num_samples, batch_size):
    # Reorder sample indices
    reordered_sample_indices = np.random.choice(num_samples, num_samples, replace
    # Generate batch indices for batch processing
    batch_indices = np.split(reordered_sample_indices, np.arange(batch_size, len(r
    return(batch_indices)
```

---

Train the 1-hidden layer neural network (128 nodes) using batch training with batch size = 100

---

```
In [ ]: ## Train the 1-hidden layer neural network (128 nodes)
## using batch training with batch size = 100
learning_rate = 1e-3 # learning rate
batch_size = 100 # batch size
nepochs = 200 # number of epochs
loss_epoch = np.empty(nepochs, dtype = np.float64) # create empty array to store

# Neural network architecture

dlayer1 = Dense(num_features, 128) # define dense layer 1
```

```

alayer1 = ReLU() # ReLU activation Layer 1
dlayer2 = Dense(128, num_labels) # define dense Layer 2
softmax = Softmax() # define softmax activation layer

# Steps: run over each sample in the batch, calculate loss, gradient of loss,
# and update weights.

epoch = 0
while epoch < nepochs:
    batch_indices = generate_batch_indices(num_samples, batch_size)
    loss = 0
    for b in range(len(batch_indices)):
        dlayer1.forward(X_train[:, batch_indices[b]]) # forward prop dense Layer 1 w
        alayer1.forward(dlayer1.output) # forward prop activation layer 1
        dlayer2.forward(alayer1.output) # forward prop dense Layer 2
        softmax.forward(dlayer2.output) # Softmax activate
        loss += cce(Y_train[:, batch_indices[b]], softmax.output) # calculate loss
        # Backward prop starts here
        grad = cce_gradient(Y_train[:, batch_indices[b]], softmax.output)
        grad = softmax.backward(grad)
        grad = dlayer2.backward(grad, learning_rate)
        grad = alayer1.backward(grad)
        grad = dlayer1.backward(grad, learning_rate)
    loss_epoch[epoch] = loss/len(batch_indices)
    print('Epoch %d: loss = %f'%(epoch+1, loss_epoch[epoch]))
    epoch = epoch + 1

```

Epoch 1: loss = 2.298796  
Epoch 2: loss = 2.290666  
Epoch 3: loss = 2.278866  
Epoch 4: loss = 2.260469  
Epoch 5: loss = 2.231436  
Epoch 6: loss = 2.186450  
Epoch 7: loss = 2.119613  
Epoch 8: loss = 2.026019  
Epoch 9: loss = 1.904079  
Epoch 10: loss = 1.758383  
Epoch 11: loss = 1.600563  
Epoch 12: loss = 1.445367  
Epoch 13: loss = 1.304079  
Epoch 14: loss = 1.181739  
Epoch 15: loss = 1.078536  
Epoch 16: loss = 0.992365  
Epoch 17: loss = 0.920443  
Epoch 18: loss = 0.860121  
Epoch 19: loss = 0.809190  
Epoch 20: loss = 0.765806  
Epoch 21: loss = 0.728518  
Epoch 22: loss = 0.696218  
Epoch 23: loss = 0.667946  
Epoch 24: loss = 0.643071  
Epoch 25: loss = 0.620950  
Epoch 26: loss = 0.601229  
Epoch 27: loss = 0.583466  
Epoch 28: loss = 0.567419  
Epoch 29: loss = 0.552848  
Epoch 30: loss = 0.539558  
Epoch 31: loss = 0.527372  
Epoch 32: loss = 0.516182  
Epoch 33: loss = 0.505854  
Epoch 34: loss = 0.496290  
Epoch 35: loss = 0.487452  
Epoch 36: loss = 0.479217  
Epoch 37: loss = 0.471535  
Epoch 38: loss = 0.464388  
Epoch 39: loss = 0.457676  
Epoch 40: loss = 0.451380  
Epoch 41: loss = 0.445471  
Epoch 42: loss = 0.439917  
Epoch 43: loss = 0.434687  
Epoch 44: loss = 0.429714  
Epoch 45: loss = 0.425027  
Epoch 46: loss = 0.420613  
Epoch 47: loss = 0.416385  
Epoch 48: loss = 0.412408  
Epoch 49: loss = 0.408594  
Epoch 50: loss = 0.404990  
Epoch 51: loss = 0.401515  
Epoch 52: loss = 0.398215  
Epoch 53: loss = 0.395049  
Epoch 54: loss = 0.392032  
Epoch 55: loss = 0.389137  
Epoch 56: loss = 0.386319  
Epoch 57: loss = 0.383674  
Epoch 58: loss = 0.381091  
Epoch 59: loss = 0.378608  
Epoch 60: loss = 0.376200

Epoch 61: loss = 0.373904  
Epoch 62: loss = 0.371687  
Epoch 63: loss = 0.369507  
Epoch 64: loss = 0.367431  
Epoch 65: loss = 0.365408  
Epoch 66: loss = 0.363457  
Epoch 67: loss = 0.361545  
Epoch 68: loss = 0.359706  
Epoch 69: loss = 0.357909  
Epoch 70: loss = 0.356151  
Epoch 71: loss = 0.354478  
Epoch 72: loss = 0.352796  
Epoch 73: loss = 0.351214  
Epoch 74: loss = 0.349646  
Epoch 75: loss = 0.348095  
Epoch 76: loss = 0.346602  
Epoch 77: loss = 0.345152  
Epoch 78: loss = 0.343730  
Epoch 79: loss = 0.342323  
Epoch 80: loss = 0.340955  
Epoch 81: loss = 0.339630  
Epoch 82: loss = 0.338322  
Epoch 83: loss = 0.337028  
Epoch 84: loss = 0.335787  
Epoch 85: loss = 0.334531  
Epoch 86: loss = 0.333314  
Epoch 87: loss = 0.332135  
Epoch 88: loss = 0.330967  
Epoch 89: loss = 0.329803  
Epoch 90: loss = 0.328660  
Epoch 91: loss = 0.327571  
Epoch 92: loss = 0.326490  
Epoch 93: loss = 0.325395  
Epoch 94: loss = 0.324328  
Epoch 95: loss = 0.323273  
Epoch 96: loss = 0.322270  
Epoch 97: loss = 0.321230  
Epoch 98: loss = 0.320227  
Epoch 99: loss = 0.319253  
Epoch 100: loss = 0.318266  
Epoch 101: loss = 0.317312  
Epoch 102: loss = 0.316366  
Epoch 103: loss = 0.315414  
Epoch 104: loss = 0.314482  
Epoch 105: loss = 0.313578  
Epoch 106: loss = 0.312667  
Epoch 107: loss = 0.311766  
Epoch 108: loss = 0.310879  
Epoch 109: loss = 0.310000  
Epoch 110: loss = 0.309133  
Epoch 111: loss = 0.308274  
Epoch 112: loss = 0.307422  
Epoch 113: loss = 0.306571  
Epoch 114: loss = 0.305750  
Epoch 115: loss = 0.304939  
Epoch 116: loss = 0.304109  
Epoch 117: loss = 0.303307  
Epoch 118: loss = 0.302498  
Epoch 119: loss = 0.301716  
Epoch 120: loss = 0.300912

Epoch 121: loss = 0.300138  
Epoch 122: loss = 0.299354  
Epoch 123: loss = 0.298585  
Epoch 124: loss = 0.297826  
Epoch 125: loss = 0.297065  
Epoch 126: loss = 0.296320  
Epoch 127: loss = 0.295556  
Epoch 128: loss = 0.294837  
Epoch 129: loss = 0.294093  
Epoch 130: loss = 0.293351  
Epoch 131: loss = 0.292623  
Epoch 132: loss = 0.291910  
Epoch 133: loss = 0.291186  
Epoch 134: loss = 0.290481  
Epoch 135: loss = 0.289774  
Epoch 136: loss = 0.289066  
Epoch 137: loss = 0.288379  
Epoch 138: loss = 0.287672  
Epoch 139: loss = 0.286978  
Epoch 140: loss = 0.286311  
Epoch 141: loss = 0.285612  
Epoch 142: loss = 0.284897  
Epoch 143: loss = 0.284253  
Epoch 144: loss = 0.283598  
Epoch 145: loss = 0.282926  
Epoch 146: loss = 0.282264  
Epoch 147: loss = 0.281599  
Epoch 148: loss = 0.280949  
Epoch 149: loss = 0.280286  
Epoch 150: loss = 0.279642  
Epoch 151: loss = 0.278997  
Epoch 152: loss = 0.278335  
Epoch 153: loss = 0.277712  
Epoch 154: loss = 0.277075  
Epoch 155: loss = 0.276441  
Epoch 156: loss = 0.275797  
Epoch 157: loss = 0.275193  
Epoch 158: loss = 0.274570  
Epoch 159: loss = 0.273931  
Epoch 160: loss = 0.273330  
Epoch 161: loss = 0.272713  
Epoch 162: loss = 0.272083  
Epoch 163: loss = 0.271514  
Epoch 164: loss = 0.270899  
Epoch 165: loss = 0.270297  
Epoch 166: loss = 0.269698  
Epoch 167: loss = 0.269099  
Epoch 168: loss = 0.268503  
Epoch 169: loss = 0.267937  
Epoch 170: loss = 0.267346  
Epoch 171: loss = 0.266752  
Epoch 172: loss = 0.266177  
Epoch 173: loss = 0.265606  
Epoch 174: loss = 0.265012  
Epoch 175: loss = 0.264461  
Epoch 176: loss = 0.263890  
Epoch 177: loss = 0.263322  
Epoch 178: loss = 0.262761  
Epoch 179: loss = 0.262185  
Epoch 180: loss = 0.261618




Epoch 181: loss = 0.261066  
Epoch 182: loss = 0.260530  
Epoch 183: loss = 0.259976  
Epoch 184: loss = 0.259423  
Epoch 185: loss = 0.258878  
Epoch 186: loss = 0.258333  
Epoch 187: loss = 0.257789  
Epoch 188: loss = 0.257250  
Epoch 189: loss = 0.256730  
Epoch 190: loss = 0.256184  
Epoch 191: loss = 0.255640  
Epoch 192: loss = 0.255129  
Epoch 193: loss = 0.254594  
Epoch 194: loss = 0.254066  
Epoch 195: loss = 0.253542  
Epoch 196: loss = 0.253010  
Epoch 197: loss = 0.252497  
Epoch 198: loss = 0.251966  
Epoch 199: loss = 0.251462  
Epoch 200: loss = 0.250942

---

Plot training loss vs. epoch

---

```
In [ ]: # Plot training loss as a function of epoch:  
plt.plot(loss_epoch)  
plt.xlabel('Epoch')  
plt.ylabel('Loss value')  
plt.show()
```



Test performance on test data

---

```
In [ ]: dlayer1.forward(X_test)
        alayer1.forward(dlayer1.output)
        dlayer2.forward(alayer1.output)
        softmax.forward(dlayer2.output)
        ypred = np.argmax(softmax.output.T, axis = 1)
        print(ypred)
        ytrue = np.argmax(Y_test.T, axis = 1)
        print(ytrue)
        np.mean(ytrue == ypred)
```

```
[7 2 1 ... 4 5 6]
[7 2 1 ... 4 5 6]
```

```
Out[ ]: 0.9305
```